

Ask the Fed

Retrieval evaluation over FOMC minutes and statements

Generated 2026-09-03 · 10,952 passages · 142 meetings · 2009-01-28 to 2026-07-29

What was measured

FOMC minutes are formulaic: the same sections say nearly the same things, meeting after meeting, for years. Semantic search alone therefore returns the right sentence from the wrong year, and that is the dominant failure mode of this corpus.

The system's defence is to parse the constraints a question already states - its period, and the document it names - and apply them in SQL *before* the vector search runs. The table below is the same 22 questions run twice against the same store: once searching everything, once narrowing first.

Result

METRIC	SEARCH EVERYTHING	NARROW FIRST	CHANGE
Passages from the wrong meeting	72.2%	3.5%	-68.7 pts
Gold passage retrieved (recall@8)	86.4%	100.0%	+13.6 pts
Some passage from the right meeting	90.9%	100.0%	+9.1 pts
Mean reciprocal rank	0.441	0.915	+0.474
Chunks searched per question	10,952	98	-10,854

The first row is the finding. Without metadata filtering, 72.2% of retrieved passages come from a meeting other than the one asked about. Filtering first brings that to 3.5%.

Where the correct passage ranked (narrow first, k=8)

POSITION	QUESTIONS
Ranked first	19
Ranked second	2
Ranked third	0
Rank 4 or worse	1
Not retrieved at all	0

Corpus

YEAR	2009	2010	2011	2012	2013	2014	2015	2016	2017	2018	2019	2020	2021	2022	2023	2024	2025	2026
Meetings	8	8	8	8	8	8	8	8	8	8	8	9	8	8	8	8	8	5

Passages by section

SECTION	PASSAGES
Committee Policy Action	1,997
Participants' Views on Current Conditions and the Economic Outlook	1,775
Staff Review of the Financial Situation	1,605
Staff Review of the Economic Situation	1,289
Preamble	1,216
Policy Statement	790
Developments in Financial Markets and Open Market Operations	761
Annual Organizational Matters	474

Documents: {"minutes": 10162, "statement": 790}. Every passage carries its meeting date and section, and the section label is what identifies whose view it is - the staff's, the participants', or the Committee's.

Limits of this evaluation

- **The gold set covers 11 of 18 years.** No pairs exist for 2009, 2010, 2011, 2012, 2013, 2014, 2015. This report therefore shows that expanding the corpus did not damage retrieval on the years it does cover - not that retrieval works on the years it does not.
- **Every gold question names its period.** So this measures "when a question states a period, does filtering help?", not "does filtering help on arbitrary questions". Questions with no temporal cue fall back to searching the whole corpus and are unrepresented here.
- **22 pairs is a small sample.** The direction is solid; the third decimal place is noise. Any change worth less than roughly two questions cannot be distinguished from chance at this size.
- **Answer generation is switched off** and therefore untested. The domain rules it must follow - preserve Fed quantifiers verbatim, never attribute a staff forecast to the Committee - are enforced only by a prompt that has not been run.